

Module 5

Discipline for Trustworthy AI Output

50 minutes

The hour

- 5 min — framing
- 15 min — three concepts (plan-first, golden rule, bidirectional sync)
- 25 min — review a PR against its plan
- 5 min — debrief

10-minute break after.

The gap

You already write a lot of the code using AI.

You already do AI-assisted PR review.

But few teams have a discipline for drift.

AI generates code. The code drifts from intent.

Your AI reviewer says "looks fine."

AI confidently reviewing AI confidently wrong.

AI speed isn't a contest of model IQ —

it's a contest of how clearly you can frame what you want.

Plan-first

Before code, write down what you're trying to do.

- Acceptance criteria
- Constraints
- Out of scope

The plan is the spec. Everything you generate is measured against it.

Defends against: **discarded specs, one-shot generation.**

The golden rule

When reality diverges from the plan,
fix the prompt first.

- If the new reality is right → the plan was wrong. Update the plan.
- If the new reality is wrong → fix the code.
- The wrong move is doing **nothing**. Drift compounds.

Defends against: **code-first divergence**.

Bidirectional sync

Two legitimate update flows:

- **Logic correction** (behavior change) → prompt-first
- **Refactor** (structure change) → code-first, sync back

Anything else is drift.

Defends against: **wholesale regeneration**.

The exercise

A teammate opened a PR for `list_routes` .
You all planned it together last sprint.

Read through these:

- `modules/module-5/sample/plan.md`
- `modules/module-5/sample/pr-description.md`
- `modules/module-5/sample/proposed/src/tools/list-routes.ts`

Would you approve it?

Reviewing — what to do

For each divergence you find:

1. **Identify** it precisely. Which line? Which acceptance criterion?
2. **Decide** who wins — the plan or the code? Why?
3. **Decide** what action follows — change the code, or update the plan?

Push back on Claude's first answer. Walk every acceptance criterion.

Coming up: Module 6

Discipline scales into automation.

AI agents in pipelines that triage failures before a human opens the logs.

10-minute break.