Module 6

AI in CI/CD on Azure DevOps

50 minutes · last module

The hour

- 10 min — honest expectations
- 20 min — live demo: build-failure triage
- 15 min — variations + "what would you wire up first?"
- 5 min — handoff

Three honest claims

1. Microsoft put their AI weight on GitHub.

ADO has MCP, Boards Copilot, YAML authoring. No pipeline-runtime AI.

2. ADO has primitives, not products.

Pipeline step. REST API. Key Vault. You wire them.

3. Cost is your problem.

Selective triggers. Token caps. Model tiering. Prompt caching.

The pattern

A pipeline step that calls an LLM.

```
read secret  →  gather context  →  call LLM  →  post somewhere
```

Three composable pieces:

- Variable group + Key Vault
- REST callback (`$(System.AccessToken)`)
- Prompt file

Same pattern, different triggers, different outputs.

The prompt is the product

The script is plumbing.

Swap the prompt → different behavior, same pipeline.

Vague prompts → the agent guesses, hallucinates, or misses what matters.

Highest-leverage line per step. Iterate it.

What would you wire up first?

What's the pipeline step that costs you mental energy every week?

- Logs you skim
- Output you summarize
- Releases you describe
- Test scaffolds you write

What would you NOT want AI to touch?

Adapting it to your repo

`participant-materials/module-6-handoff-checklist.md`

1. Variable group + Key Vault
2. Pipeline YAML — three places to edit
3. Prompt — the leverage
4. Cost cap — `max_tokens`, timeout, selective triggers
5. Try it

Slack me when you ship one. Two-week window.

Thanks

Bundle within 24h:

- Workshop repo (FleetOps sample, MCP, plans, skills)
- This ADO repo + handoff checklist
- Takeaways doc
- SPDD reference (Module 5)
- Research record (Module 6)

Ship one thing.